

# Image Caption Generator with CNN & LSTM

Python based project where we will use deep learning techniques of Convolutional Neural Networks and a type of Recurrent Neural Network (LSTM) together.

```
In [1]: import string
import numpy as np
from PIL import Image
import os
from pickle import dump, load
import numpy as np

from keras.applications.xception import Xception, preprocess_input
from keras.preprocessing.image import load_img, img_to_array
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from keras.utils import to_categorical
from keras.layers import add
from keras.models import Model, load_model
from keras.layers import Input, Dense, LSTM, Embedding, Dropout

# small library for seeing the progress of loops.
from tqdm.notebook import tqdm as tqdm
tqdm().pandas()

0it [00:00, ?it/s]
```

```
In [2]: # Loading a text file into memory
def load_doc(filename):
    # Opening the file as read only
    file = open(filename, 'r')
    text = file.read()
    file.close()
    return text
```

```
In [3]: # get all imgs with their captions
def all_img_captions(filename):
```

```

file = load_doc(filename)
captions = file.split('\n')
descriptions = {}
for caption in captions[:-1]:
    img, caption = caption.split('\t')
    if img[:-2] not in descriptions:
        descriptions[img[:-2]] = [caption]
    else:
        descriptions[img[:-2]].append(caption)
return descriptions

```

In [4]: *##Data cleaning- lower casing, removing punctuations and words containing numbers*

```

def cleaning_text(captions):
    table = str.maketrans('', '', string.punctuation)
    for img, caps in captions.items():
        for i, img_caption in enumerate(caps):

            img_caption.replace("-", " ")
            desc = img_caption.split()

            #converts to lower case
            desc = [word.lower() for word in desc]
            #remove punctuation from each token
            desc = [word.translate(table) for word in desc]
            #remove hanging 's and a
            desc = [word for word in desc if (len(word)>1)]
            #remove tokens with numbers in them
            desc = [word for word in desc if (word.isalpha())]
            #convert back to string

            img_caption = ' '.join(desc)
            captions[img][i] = img_caption
    return captions

```

In [5]: **def** text\_vocabulary(descriptions):  
*# build vocabulary of all unique words*  
 vocab = set()

```

for key in descriptions.keys():
    [vocab.update(d.split()) for d in descriptions[key]]

```

```
return vocab
```

```
In [6]: #All descriptions in one file
def save_descriptions(descriptions, filename):
    lines = list()
    for key, desc_list in descriptions.items():
        for desc in desc_list:
            lines.append(key + '\t' + desc )
    data = "\n".join(lines)
    file = open(filename,"w")
    file.write(data)
    file.close()
```

```
In [7]: # Set these path according to project folder in you system
dataset_text = "C:/Users/vidit/Desktop/MS in DS/Projects/Image Caption Generator/Flickr8k_text"
dataset_images = "C:/Users/vidit/Desktop/MS in DS/Projects/Image Caption Generator/Flickr8k_Dataset"
```

```
In [8]: #we prepare our text data
filename = dataset_text + "/" + "Flickr8k.token.txt"
#Loading the file that contains all data
#mapping them into descriptions dictionary img to 5 captions
descriptions = all_img_captions(filename)
print("Length of descriptions =" ,len(descriptions))

Length of descriptions = 8092
```

```
In [9]: #cleaning the descriptions
clean_descriptions = cleaning_text(descriptions)
```

```
In [10]: #building vocabulary
vocabulary = text_vocabulary(clean_descriptions)
print("Length of vocabulary = ", len(vocabulary))

Length of vocabulary = 8763
```

```
In [11]: #saving each description to file
save_descriptions(clean_descriptions, "descriptions.txt")
```

```
In [12]: def extract_features(directory):
        model = Xception( include_top=False, pooling='avg' )
```

```

features = {}
for img in tqdm(os.listdir(directory)):
    filename = directory + "/" + img
    image = Image.open(filename)
    image = image.resize((299,299))
    image = np.expand_dims(image, axis=0)
    #image = preprocess_input(image)
    image = image/127.5
    image = image - 1.0

    feature = model.predict(image)
    features[img] = feature
return features

```

## Extracting the feature vector from all images

```

In [23]: #2048 feature vector
features = extract_features(dataset_images)
dump(features, open("features.p", "wb"))

```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/xception/xception\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/xception/xception_weights_tf_dim_ordering_tf_kernels_notop.h5)  
83683744/83683744 [=====] - 35s 0us/step  
0%| | 0/8091 [00:00<?, ?it/s]

```

1/1 [=====] - 0s 180ms/step
1/1 [=====] - 0s 186ms/step
1/1 [=====] - 0s 183ms/step
1/1 [=====] - 0s 186ms/step
1/1 [=====] - 0s 183ms/step
1/1 [=====] - 0s 183ms/step
1/1 [=====] - 0s 214ms/step
1/1 [=====] - 0s 194ms/step
1/1 [=====] - 0s 205ms/step
1/1 [=====] - 0s 198ms/step
1/1 [=====] - 0s 206ms/step
1/1 [=====] - 0s 201ms/step
1/1 [=====] - 0s 186ms/step
1/1 [=====] - 0s 190ms/step
1/1 [=====] - 0s 185ms/step
1/1 [=====] - 0s 192ms/step
1/1 [=====] - 0s 179ms/step
1/1 [=====] - 0s 181ms/step
1/1 [=====] - 0s 188ms/step
1/1 [=====] - 0s 180ms/step
1/1 [=====] - 0s 181ms/step
1/1 [=====] - 0s 183ms/step
1/1 [=====] - 0s 190ms/step
1/1 [=====] - 0s 185ms/step
1/1 [=====] - 0s 193ms/step
1/1 [=====] - 0s 218ms/step
1/1 [=====] - 0s 197ms/step

```

```
In [13]: features = load(open("features.p", "rb"))
```

```
In [14]: #Load the data
def load_photos(filename):
    file = load_doc(filename)
    photos = file.split("\n")[:-1]
    return photos
```

```
In [15]: def load_clean_descriptions(filename, photos):
    #Loading clean_descriptions
    file = load_doc(filename)
    descriptions = {}
    for line in file.split("\n"):
```

```

words = line.split()
if len(words)<1 :
    continue

image, image_caption = words[0], words[1:]

if image in photos:
    if image not in descriptions:
        descriptions[image] = []
    desc = '<start> ' + " ".join(image_caption) + ' <end>'
    descriptions[image].append(desc)

return descriptions

```

```

In [16]: def load_features(photos):
         #loading all features
         all_features = load(open("features.p", "rb"))
         #selecting only needed features
         features = {k:all_features[k] for k in photos}
         return features

```

```

In [17]: filename = dataset_text + "/" + "Flickr_8k.trainImages.txt"

         #train = loading_data(filename)
         train_imgs = load_photos(filename)
         train_descriptions = load_clean_descriptions("descriptions.txt", train_imgs)
         train_features = load_features(train_imgs)

```

## Tokenizing the vocabulary

```

In [18]: #converting dictionary to clean list of descriptions
         def dict_to_list(descriptions):
             all_desc = []
             for key in descriptions.keys():
                 [all_desc.append(d) for d in descriptions[key]]
             return all_desc

```

```
In [19]: #creating tokenizer class  
#this will vectorise text corpus  
#each integer will represent token in dictionary
```

```
from keras.preprocessing.text import Tokenizer  
  
def create_tokenizer(descriptions):  
    desc_list = dict_to_list(descriptions)  
    tokenizer = Tokenizer()  
    tokenizer.fit_on_texts(desc_list)  
    return tokenizer
```

```
In [20]: # give each word a index, and store that into tokenizer.p pickle file  
tokenizer = create_tokenizer(train_descriptions)  
dump(tokenizer, open('tokenizer.p', 'wb'))  
vocab_size = len(tokenizer.word_index) + 1  
vocab_size
```

Out[20]: 7577

```
In [21]: #calculate maximum length of descriptions  
def max_length(descriptions):  
    desc_list = dict_to_list(descriptions)  
    return max(len(d.split()) for d in desc_list)  
  
max_length = max_length(descriptions)  
max_length
```

Out[21]: 32

```
In [22]: features['1000268201_693b08cb0e.jpg'][0]
```

```
Out[22]: array([0.47340965, 0.01730897, 0.07334232, ..., 0.08557969, 0.02102296,  
                0.23765543], dtype=float32)
```

## Define the model

- Photo feature extractor - we extracted features from pretrained model Xception.
- Sequence processor - word embedding layer that handles text, followed by LSTM

- Decoder - Both 1 and 2 model produce fixed length vector. They are merged together and processed by dense layer to make final prediction

```
In [23]: #create input-output sequence pairs from the image description.

#data generator, used by model.fit_generator()
def data_generator(descriptions, features, tokenizer, max_length):
    while 1:
        for key, description_list in descriptions.items():
            #retrieve photo features
            feature = features[key][0]
            input_image, input_sequence, output_word = create_sequences(tokenizer, max_length, description_list, feature)
            yield [[input_image, input_sequence], output_word]

def create_sequences(tokenizer, max_length, desc_list, feature):
    X1, X2, y = list(), list(), list()
    # walk through each description for the image
    for desc in desc_list:
        # encode the sequence
        seq = tokenizer.texts_to_sequences([desc])[0]
        # split one sequence into multiple X,y pairs
        for i in range(1, len(seq)):
            # split into input and output pair
            in_seq, out_seq = seq[:i], seq[i]
            # pad input sequence
            in_seq = pad_sequences([in_seq], maxlen=max_length)[0]
            # encode output sequence
            out_seq = to_categorical([out_seq], num_classes=vocab_size)[0]
            # store
            X1.append(feature)
            X2.append(in_seq)
            y.append(out_seq)
    return np.array(X1), np.array(X2), np.array(y)
```

```
In [24]: [a,b],c = next(data_generator(train_descriptions, features, tokenizer, max_length))
a.shape, b.shape, c.shape
```

```
Out[24]: ((47, 2048), (47, 32), (47, 7577))
```



## Defining the CNN-RNN model

```
In [25]: from keras.utils import plot_model

# define the captioning model
def define_model(vocab_size, max_length):

    # features from the CNN model squeezed from 2048 to 256 nodes
    inputs1 = Input(shape=(2048,))
    fe1 = Dropout(0.5)(inputs1)
    fe2 = Dense(256, activation='relu')(fe1)

    # LSTM sequence model
    inputs2 = Input(shape=(max_length,))
    se1 = Embedding(vocab_size, 256, mask_zero=True)(inputs2)
    se2 = Dropout(0.5)(se1)
    se3 = LSTM(256)(se2)

    # Merging both models
    decoder1 = add([fe2, se3])
    decoder2 = Dense(256, activation='relu')(decoder1)
    outputs = Dense(vocab_size, activation='softmax')(decoder2)

    # tie it together [image, seq] [word]
    model = Model(inputs=[inputs1, inputs2], outputs=outputs)
    model.compile(loss='categorical_crossentropy', optimizer='adam')

    # summarize model
    print(model.summary())
    plot_model(model, to_file='model.png', show_shapes=True)

    return model
```

## Training the model

```
In [27]: # train our model
print('Dataset: ', len(train_imgs))
print('Descriptions: train=', len(train_descriptions))
```

```
print('Photos: train=', len(train_features))
print('Vocabulary Size:', vocab_size)
print('Description Length: ', max_length)

model = define_model(vocab_size, max_length)
epochs = 10
steps = len(train_descriptions)
# making a directory models to save our models
os.mkdir("models")
for i in range(epochs):
    generator = data_generator(train_descriptions, train_features, tokenizer, max_length)
    model.fit_generator(generator, epochs=1, steps_per_epoch= steps, verbose=1)
    model.save("models/model_" + str(i) + ".h5")
```

Dataset: 6000  
 Descriptions: train= 6000  
 Photos: train= 6000  
 Vocabulary Size: 7577  
 Description Length: 32  
 Model: "model\_1"

Layer (type)	Output Shape	Param #	Connected to
input_4 (InputLayer)	[(None, 32)]	0	[]
input_3 (InputLayer)	[(None, 2048)]	0	[]
embedding_1 (Embedding)	(None, 32, 256)	1939712	['input_4[0][0]']
dropout_2 (Dropout)	(None, 2048)	0	['input_3[0][0]']
dropout_3 (Dropout)	(None, 32, 256)	0	['embedding_1[0][0]']
dense_3 (Dense)	(None, 256)	524544	['dropout_2[0][0]']
lstm_1 (LSTM)	(None, 256)	525312	['dropout_3[0][0]']
add_1 (Add)	(None, 256)	0	['dense_3[0][0]', 'lstm_1[0][0]']
dense_4 (Dense)	(None, 256)	65792	['add_1[0][0]']
dense_5 (Dense)	(None, 7577)	1947289	['dense_4[0][0]']

=====  
 Total params: 5002649 (19.08 MB)  
 Trainable params: 5002649 (19.08 MB)  
 Non-trainable params: 0 (0.00 Byte)

None

You must install pydot (`pip install pydot`) and install graphviz (see instructions at <https://graphviz.gitlab.io/download/>) for plot\_model to work.

C:\Users\vidit\AppData\Local\Temp\ipykernel\_17764\2846483122.py:15: UserWarning: `Model.fit\_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.

model.fit\_generator(generator, epochs=1, steps\_per\_epoch= steps, verbose=1)

6000/6000 [=====] - 1446s 240ms/step - loss: 4.5192

```
C:\Users\vidit\AppData\Roaming\Python\Python310\site-packages\keras\src\engine\training.py:3079: UserWarning: You are saving your model as an HDF5 file via `model.save()`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')`.
  saving_api.save_model(
```

```
6000/6000 [=====] - 1444s 241ms/step - loss: 3.6668
6000/6000 [=====] - 1800s 300ms/step - loss: 3.3763
6000/6000 [=====] - 1778s 296ms/step - loss: 3.2024
6000/6000 [=====] - 1868s 311ms/step - loss: 3.0830
6000/6000 [=====] - 1903s 317ms/step - loss: 2.9933
6000/6000 [=====] - 1878s 313ms/step - loss: 2.9248
6000/6000 [=====] - 1884s 314ms/step - loss: 2.8677
6000/6000 [=====] - 1915s 319ms/step - loss: 2.8221
6000/6000 [=====] - 1963s 327ms/step - loss: 2.7828
```

## Testing the model

```
In [28]: from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from keras.applications.xception import Xception
from keras.models import load_model
from pickle import load
import numpy as np
from PIL import Image
import matplotlib.pyplot as plt
import argparse
```

```
In [59]: img_path = "C:/Users/vidit/Desktop/MS in DS/Projects/Image Caption Generator/Flicker8k_Dataset/3665179773_dd217416fc.
#img_path = "C:/Users/vidit/Desktop/MS in DS/Projects/Image Caption Generator/Flicker8k_Dataset/3471463779_64084b686c
#img_path = "C:/Users/vidit/Desktop/MS in DS/Projects/Image Caption Generator/Flicker8k_Dataset/3457572788_e1fe4f6486
#img_path = "C:/Users/vidit/Desktop/MS in DS/Projects/Image Caption Generator/Flicker8k_Dataset/421808539_57abee6d55.
#img_path = "C:/Users/vidit/Desktop/MS in DS/Projects/Image Caption Generator/Flicker8k_Dataset/242109387_e497277e07.
```

```
In [70]: def result(img_path):
    def extract_features(filename, model):
        try:
            image = Image.open(filename)

        except:
            print("ERROR: Couldn't open image! Make sure the image path and extension is correct")
            image = image.resize((299,299))
```

```

image = np.array(image)
# for images that has 4 channels, we convert them into 3 channels
if image.shape[2] == 4:
    image = image[..., :3]
image = np.expand_dims(image, axis=0)
image = image/127.5
image = image - 1.0
feature = model.predict(image)
return feature

def word_for_id(integer, tokenizer):
    for word, index in tokenizer.word_index.items():
        if index == integer:
            return word
    return None

def generate_desc(model, tokenizer, photo, max_length):
    in_text = 'start'
    for i in range(max_length):
        sequence = tokenizer.texts_to_sequences([in_text])[0]
        sequence = pad_sequences([sequence], maxlen=max_length)
        pred = model.predict([photo, sequence], verbose=0)
        pred = np.argmax(pred)
        word = word_for_id(pred, tokenizer)
        if word is None:
            break
        in_text += ' ' + word
        if word == 'end':
            break
    return in_text

max_length = 32
tokenizer = load(open("tokenizer.p", "rb"))
model = load_model('models/model_9.h5')
xception_model = Xception(include_top=False, pooling="avg")

photo = extract_features(img_path, xception_model)
img = Image.open(img_path)

```

```

description = generate_desc(model, tokenizer, photo, max_length)
print("\n\n")
print(description)
plt.imshow(img)
plt.show()

```

```


In [71]: Arr=["421808539_57abee6d55", "242109387_e497277e07", "3457572788_e1fe4f6480", "3665179773_dd217416fc", "3471463779_64084k

for i in Arr:
    img_path = "C:/Users/vidit/Desktop/MS in DS/Projects/Image Caption Generator/Flicker8k_Dataset/"+i+".jpg"
    result(img_path)

```

1/1 [=====] - 1s 953ms/step

start dog is running through the snow end




1/1 [=====] - 1s 823ms/step

start little boy in pink shirt plays on playground end



1/1 [=====] - 1s 784ms/step

start skier is skiing down snowy hill end



1/1 [=====] - 1s 1s/step

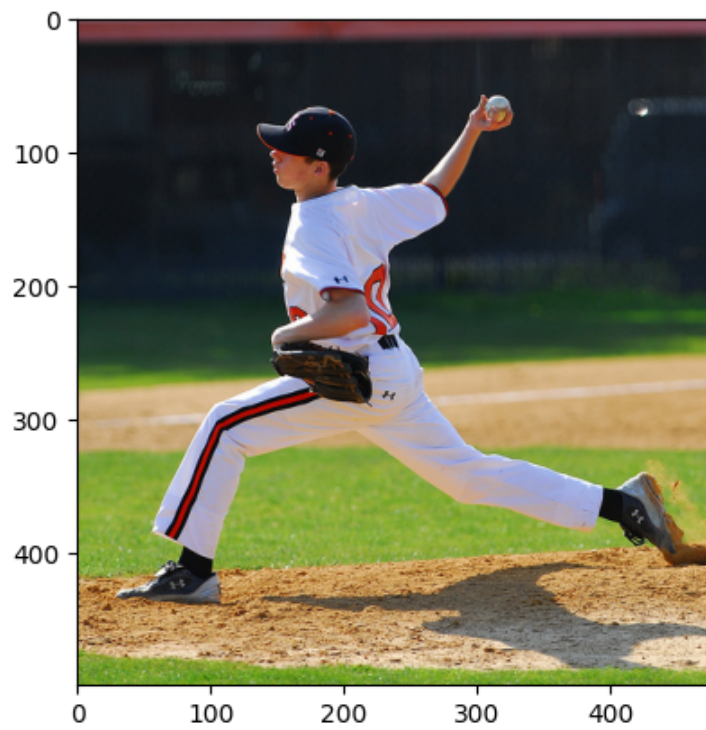
start man in red shirt is sitting on bench with dog end





1/1 [=====] - 1s 878ms/step

start baseball player in white shirt is throwing ball end



In [ ]:

In [ ]:

In [ ]: